

# Linear Regression

Machine Learning Practice

Dr. Ashish Tendulkar

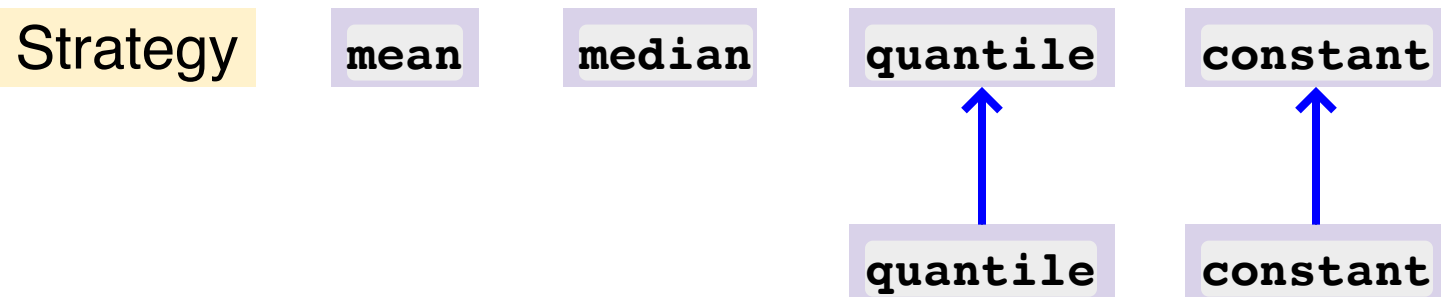
IIT Madras

# How to build baseline regression model?

**DummyRegressor** helps in creating a **baseline** for regression.

```
1 from sklearn.dummy import DummyRegressor
2
3 dummy_regr = DummyRegressor(strategy="mean")
4 dummy_regr.fit(X_train, y_train)
5 dummy_regr.predict(X_test)
6 dummy_regr.score(X_test, y_test)
```

- It makes a prediction as specified by the **strategy**.
- Strategy is based on **some statistical property** of the training set or **user specified value**.



# How is **Linear Regression** model trained?

**Step 1:** Instantiate **object** of a suitable **linear regression estimator** from one of the following two options

Normal  
equation

```
1 from sklearn.linear_model import LinearRegression
2 linear_regressor = LinearRegression()
```

Iterative  
optimization

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor()
```

**Step 2:** Call **fit** method on **linear regression object** with **training feature matrix** and **label vector** as arguments.

```
1 # Model training with feature matrix X_train and
2 # label vector or matrix y_train
3 linear_regressor.fit(X_train, y_train)
```

Works for both single and multi-output regression.

# SGDRegressor Estimator

# SGDRegressor Estimator

- Implements **stochastic gradient descent**
- Use for large training set up (> 10k samples)
- Provides **greater control on optimization process** through provision for **hyperparameter** settings.

- `loss= 'squared error'`
- `loss = 'huber'`

- `penalty = 'l1'`
- `penalty = 'l2'`
- `penalty = 'elasticnet'`

SGDRegressor



- `learning_rate = 'constant'`
- `learning_rate = 'optimal'`
- `learning_rate = 'invscaling'`
- `learning_rate = 'adaptive'`

- `early_stopping = 'True'`
- `early_stopping = 'False'`

It's a good idea to use a **random seed** of your choice while instantiating SGDRegressor object. It helps us get **reproducible results**.

Set `random_state` to seed of your choice.

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(random_state=42)
```

**Note:** In the rest of the presentation, we won't set the random seed for sake of brevity. However while coding, always set the random seed in the constructor.

# How to perform feature scaling for SGDRegressor?

SGD is **sensitive to feature scaling**, so it is **highly recommended to scale** input feature matrix.

```
1 from sklearn.linear_model import SGDRegressor
2 from sklearn.pipeline import Pipeline
3 from sklearn.preprocessing import StandardScaler
4
5 sgd = Pipeline([
6     ('feature_scaling', StandardScaler()),
7     ('sgd_regressor', SGDRegressor())])
8
9 sgd.fit(X_train, y_train)
```

- Note
- Feature scaling **is not needed** for **word frequencies** and **indicator features** as they have intrinsic scale.
  - Features extracted using PCA should be **scaled by some constant  $c$**  such that the average L2 norm of the training data equals one.

# How to shuffle training data after each epoch in SGDRegressor?

```
1 from sklearn.linear_model import SGDRegressor  
2 linear_regressor = SGDRegressor(shuffle=True)
```



# How to use set **learning rate** in SGDRegressor?

- `learning_rate = 'constant'`

- `learning_rate = 'invscaling'`

- `learning_rate = 'adaptive'`

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(random_state=42)
```

What is the default setting?

- `learning_rate = 'invscaling'`
- `eta0 = 1e-2`
- `power_t = 0.25`

Learning rate **reduces** after **every iteration**:

$$\text{eta} = \text{eta0} / \text{pow}(t, \text{power\_t})$$

**Note:** You can make changes to these parameters to speed up or slow down the training process.

## How to use set constant learning rate ?

- `learning_rate = 'constant'`

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(learning_rate='constant',
3                                eta0=1e-2)
```

Constant learning rate `eta0 = 1e-2` is used throughout the training.

## How to set adaptive learning rate?

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(learning_rate='adaptive',
3                                 eta0=1e-2)
```

- The learning rate is kept to **initial value** as long as the training loss **decreases**.
- When the **stopping criterion** is reached, the **learning rate** is divided **by 5**, and the **training loop continues**.
- The algorithm stops when the **learning rate goes below  $10^{-6}$** .

# How to set `#epochs` in `SGDRegressor`?

Set `max_iter` to desired `#epochs`. The default value is 1000.

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(max_iter=100)
```

Remember `one epoch` is `one full pass over the training data`.

## Practical tip

SGD converges after observing approximately  $10^6$  training samples. Thus, a reasonable first guess for the number of iterations for  $n$  sampled training set is

$$\text{max\_iter} = \text{np.ceil}(10^6 / n)$$

# How to set **stopping criteria** in SGDRegressor?

Option #1 `tol`, `n_iter_no_change`, `max_iter`.

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(loss='squared_error',
3                                 max_iter=500,
4                                 tol=1e-3,
5                                 n_iter_no_change=5)
```

The SGDRegressor **stops**

- when **the training loss does not improve** (`loss > best_loss - tol`) for `n_iter_no_change` consecutive epochs
- else after a maximum number of iteration `max_iter`.

# How to set **stopping criteria** in SGDRegressor?

Option #2 `early_stopping`, `validation_fraction`

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(loss='squared_error',
3                                 early_stopping=True
4                                 max_iter=500,
5                                 tol=1e-3,
6                                 validation_fraction=0.2,
7                                 n_iter_no_change=5)
```

Set aside `validation_fraction` percentage records from training set as validation set. Use `score` method to obtain validation score.

The SGDRegressor **stops** when

- **validation score** does not improve by at least `tol` for `n_iter_no_change` consecutive epochs.
- else after a maximum number of iteration `max_iter`.

# How to use different **loss functions** in SGDRegressor?

Set **loss** parameter to one of the supported values

**'squared\_error'** {studied in this course}

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(loss='squared_error')
```

It also supports other losses as documented in [sklearn API](#)

# How to use averaged SGD?

Averaged SGD updates the weight vector to **average of weights** from previous updates.

Option #1: Averaging across all updates **average=True**

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(average=True)
```



Option #2: Set **average** to int value.

Averaging begins once the total number of samples seen reaches **average**

Setting **average=10** starts averaging after seeing 10 samples

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(average=10)
```

Averaged SGD works **best** with a **larger number of features** and a **higher eta0**

# How do we initialize SGD with weight vector of the previous run?

Set `warm_start = True`

while instantiating object of `SGDRegressor`

```
1 from sklearn.linear_model import SGDRegressor
2 linear_regressor = SGDRegressor(warm_start=True)
```

By default `warm_start = False`

# How to monitor SGD loss iteration after iteration?

Make use of **warm\_start = TRUE**

```
1 sgd_reg = SGDRegressor(max_iter=1, tol=-np.infty, warm_start=True,  
2                       penalty=None, learning_rate="constant", eta0=0.0005)  
3  
4 for epoch in range(1000):  
5     sgd_reg.fit(X_train, y_train) # continues where it left off  
6     y_val_predict = sgd_reg.predict(X_val)  
7     val_error = mean_squared_error(y_val, y_val_predict)
```

# Model inspection

# How to access the weights of trained **Linear Regression** model?

$$\hat{y} = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m = \mathbf{w}^T \mathbf{x}$$

The **weights**  $w_1, w_2, \dots, w_m$  are stored in **coef\_** class variable.

```
1 linear_regressor.coef_
```

The **intercept**  $w_0$  is stored in **intercept\_** class variable.

```
1 linear_regressor.intercept_
```

**Note:** These code snippets works for both **LinearRegression** and **SGDRegressor**, and for that matter to **all regression estimators** that we will study in this module. Why?

All of them are estimators.

# Model inference

# How to make predictions on new data in **Linear Regression** model?

**Step 1:** Arrange data for prediction in a feature matrix of shape (#samples, #features) or in sparse matrix format.

**Step 2:** Call **predict** method on **linear regression object** with **feature matrix** as an argument.

```
1 # Predict labels for feature matrix X_test
2 linear_regressor.predict(X_test)
```

Same code works for **all regression estimators**.

# Model evaluation



# General steps in model evaluation

**STEP 1:** Split data into train and test

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

**STEP 2:** Fit linear regression estimator on training set.

**STEP 3:** Calculate training error (a.k.a. empirical error)

**STEP 4:** Calculate test error (a.k.a. generalization error)

Compare training and test errors

# How to evaluate trained Linear Regression model?

Using `score` method on `linear regression object`:

```
1 # Evaluation on the eval set with
2 # 1. feature matrix
3 # 2. label vector or matrix (single/multi-output)
4 linear_regressor.score(X_test, y_test)
```

The score returns  $R^2$  or coefficient of determination

residual sum of squares:

$$u = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y})$$

Sum of squared error  
(actual and predicted label)

$$R^2 = \left(1 - \frac{u}{v}\right)$$

total sum of square

Sum of squared error  
(actual and mean predicted  
label)

$$v = (\mathbf{y} - \hat{\mathbf{y}}_{\text{mean}})^T (\mathbf{y} - \hat{\mathbf{y}}_{\text{mean}})$$

The score returns  $R^2$  or coefficient of determination

$$R^2 = \left(1 - \frac{u}{v}\right)$$

When?

- The **best possible score** is **1.0**.

$u$ , sum of squared error = 0

- A **constant model** that always predicts the expected value of  $y$ , would get a **score of 0.0**.

$$u = v$$

- The **score** can be **negative** (because the model can be arbitrarily worse).

# Evaluation metrics

sklearn provides a **bunch of regression metrics** to evaluate **performance** of the **trained estimator** on the **evaluation set**.

## **mean\_absolute\_error**

```
1 from sklearn.metrics import mean_absolute_error
2 eval_score = mean_absolute_error(y_test, y_predicted)
```

## **mean\_squared\_error**

```
1 from sklearn.metrics import mean_squared_error
2 eval_score = mean_squared_error(y_test, y_predicted)
```

## **r2\_score** Same as output of **score**

```
1 from sklearn.metrics import r2_score
2 eval_score = r2_score(y_test, y_predicted)
```

These metrics can also be used in **multi-output regression** setup.

## mean\_squared\_log\_error

```
1 from sklearn.metrics import mean_squared_log_error
2 eval_score = mean_squared_log_error(y_test, y_predicted)
```

- Useful for **targets with exponential growths** like population, sales growth etc,
- **Penalizes under-estimation heavier than the over-estimation.**

## mean\_absolute\_percentage\_error

```
1 from sklearn.metrics import mean_absolute_percentage_error
2 eval_score = mean_absolute_percentage_error(y_test, y_predicted)
```

- **Sensitive to relative error.**

## median\_absolute\_error

```
1 from sklearn.metrics import median_absolute_error
2 eval_score = median_absolute_error(y_test, y_predicted)
```

- **Robust to outliers**

# How to evaluate regression model on worst case error?

Use metrics `max_error`

Worst case error on train set can be calculated as follows:

```
1 from sklearn.metrics import max_error
2 train_error = max_error(y_train, y_predicted)
```

Worst case error on test set can be calculated as follows:

```
1 from sklearn.metrics import max_error
2 test_error = max_error(y_test, y_predicted)
```

This metrics can, however, be used only for **single output regression**. It **does not support multi-output regression**.

# Scores and Errors

- Score is a metric for which higher value is better.
- Error is a metric for which lower value is better.

Convert error metric to score metric by adding **neg\_** suffix.

Function	Scoring
<code>metrics.mean_absolute_error</code>	<code>neg_mean_absolute_error</code>
<code>metrics.mean_squared_error</code>	<code>neg_mean_squared_error</code>
<code>metrics.mean_squared_error</code>	<code>neg_root_mean_squared_error</code>
<code>metrics.mean_squared_log_error</code>	<code>neg_mean_squared_log_error</code>
<code>metrics.median_absolute_error</code>	<code>neg_median_absolute_error</code>

In case, we get comparable performance on train and test with this split, is this performance guaranteed on other splits too?

- Is test set sufficiently large?
  - In case it is small, the test error obtained may be unstable and would not reflect the true test error on large test set.
- What is the chance that the easiest examples were kept aside as test by chance?
  - This if happens would lead to optimistic estimation of the true test error.

We use cross validation for robust performance evaluation.



Cross-validation performs **robust evaluation** of model performance

- by **repeated splitting** and
- providing **many training and test errors**

This enables us to **estimate variability in generalization performance** of the model.

sklearn implements the following cross validation iterators

**KFold**

**RepeatedKfold**

**LeaveOneOut**

**ShuffleSplit**

# How to obtain cross validated performance measure using **KFold**?

```
1 from sklearn.model_selection import cross_val_score
2 from sklearn.linear_model import linear_regression
3
4 lin_reg = linear_regression()
5 score = cross_val_score(lin_reg, X, y, cv=5)
```

- Uses **KFold** cross validation iterator, that divides training data into 5 folds.
- In each run, it uses **4 folds for training** and **1 for evaluation**.

## Alternate way of writing the same thing

```
1 from sklearn.model_selection import cross_val_score
2 from sklearn.model_selection import KFold
3 from sklearn.linear_model import linear_regression
4
5 lin_reg = linear_regression()
6 kfold_cv = KFold(n_splits=5, random_state=42)
7 score = cross_val_score(lin_reg, X, y, cv=kfold_cv)
```

# How to obtain cross validated performance measure using **LeaveOneOut**?

```
1 from sklearn.model_selection import cross_val_score
2 from sklearn.model_selection import LeaveOneOut
3 from sklearn.linear_model import linear_regression
4
5 lin_reg = linear_regression()
6 loocv = LeaveOneOut()
7 score = cross_val_score(lin_reg, X, y, cv=loocv)
```

which is same as

```
1 from sklearn.model_selection import cross_val_score
2 from sklearn.model_selection import KFold
3 from sklearn.linear_model import linear_regression
4
5 lin_reg = linear_regression()
6 n = X.shape[0]
7 kfold_cv = KFold(n_splits=n)
8 score = cross_val_score(lin_reg, X, y, cv=kfold_cv)
```

# How to obtain cross validated performance measure using **ShuffleSplit**?

```
1 from sklearn.linear_model import linear_regression
2 from sklearn.model_selection import cross_val_score
3 from sklearn.model_selection import ShuffleSplit
4
5 lin_reg = linear_regression()
6 shuffle_split = ShuffleSplit(n_splits=5, test_size=0.2, random_state=42)
7 score = cross_val_score(lin_reg, X, y, cv=shuffle_split)
```

It is also called **random permutation** based **cross validation strategy**.

- Generates user defined number of train/test splits.
- It is robust to class distribution.

In each iteration, it shuffles order of data samples and then splits it into train and test.

# How to specify a performance measure in **cross\_val\_score**

```
1 from sklearn.linear_model import linear_regression
2 from sklearn.model_selection import cross_val_score
3 from sklearn.model_selection import ShuffleSplit
4
5 lin_reg = linear_regression()
6 shuffle_split = ShuffleSplit(n_splits=5, test_size=0.2, random_state=42)
7 score = cross_val_score(lin_reg, X, y, cv=shuffle_split,
8                          scoring='neg_mean_absolute_error')
```

**scoring** parameter can be set to one of the scoring schemes implemented in sklearn as follows

<b>max_error</b>	<b>r2</b>
<b>neg_mean_absolute_error</b>	<b>neg_mean_squared_error</b>
<b>neg_mean_squared_log_error</b>	<b>neg_median_absolute_error</b>
<b>neg_root_mean_squared_error</b>	

# How to obtain **test scores** from different folds?

```
1 from sklearn.model_selection import cross_validate
2 from sklearn.model_selection import ShuffleSplit
3
4 cv = ShuffleSplit(n_splits=40, test_size=0.3, random_state=0)
5 cv_results = cross_validate(
6     regressor, data, target, cv=cv, scoring="neg_mean_absolute_error")
```

The results are stored in python dictionary with the following keys:

**fit\_time**

**score\_time**

**test\_score**

**estimator** (optional)

**train\_score** (optional)

# How to obtain **trained estimators** and **scores** on training data during cross validation?

- For trained estimator, set **return\_estimator = True**
- For scores on training set, set **return\_train\_score = True**

```
1 from sklearn.model_selection import cross_validate
2 from sklearn.model_selection import ShuffleSplit
3
4 cv = ShuffleSplit(n_splits=40, test_size=0.3,
5                 random_state=0)
6 cv_results = cross_validate(
7     regressor, data, target,
8     cv=cv, scoring="neg_mean_absolute_error",
9     return_train_score=True,
10    return_estimator=True)
```

The estimators can be accessed through **estimator** key of the dictionary returned by **cross\_validate**

# How to evaluate **multiple metrics** of regression in cross validation set up?

```
1 from sklearn.model_selection import cross_validate
2 from sklearn.model_selection import ShuffleSplit
3
4 cv = ShuffleSplit(n_splits=40, test_size=0.3,
5                 random_state=0)
6 cv_results = cross_validate(
7     regressor, data, target,
8     cv=cv,
9     scoring=["neg_mean_absolute_error", "neg_mean_squared_error"]
10    return_train_score=True,
11    return_estimator=True)
```

**cross\_validate** allows us to specify multiple scoring metrics  
unlike **cross\_val\_score**



# How to study effect of #samples on training and test errors?

**STEP 1:** Instantiate an object of `learning_curve` class with `estimator`, `training data`, `size`, `cross validation strategy` and `scoring scheme` as arguments.

```
1 from sklearn.model_selection import learning_curve
2
3 results = learning_curve(
4     lin_reg, X_train, y_train, train_sizes=train_sizes, cv=cv,
5     scoring="neg_mean_absolute_error")
6 train_size, train_scores, test_scores = results[:3]
7 # Convert the scores into errors
8 train_errors, test_errors = -train_scores, -test_scores
```

**STEP 2:** Plot `training and test scores` as function of the `size` of training sets. And make assessment about model fitment: `under/overfitting` or `right fit`.

# Underfitting/Overfitting diagnosis

**STEP 1:** Fit linear models with different number of features.

**STEP 2:** For each model, obtain training and test errors.

**STEP 3:** Plot #features vs error graph - one each for training and test errors.

**STEP 4:** Examine the graphs to detect under/overfitting.

We can replace #features with any other tunable hyperparameter to do this diagnosis for setting that hyperparameter to the appropriate value.