

Neural Networks

Dr. Ashish Tendulkar

IIT Madras

Machine Learning Practice

- In this week, we will study how to implement **Multilayer Perceptron** neural network models for **classification and regression** tasks with **sklearn**.

Multilayer Perceptron (MLP)

- It is a supervised learning algorithm.
- MLP learns a **non-linear function approximator** for either classification or regression depending on the given dataset.
- In **sklearn**, we implement MLP using:
 1. MLPClassifier for classification
 2. MLPRegressor for regression
- **MLPClassifier** supports **multi-class classification** by applying Softmax as the output function.
- It also supports **multi-label classification** in which a sample can belong to more than one class.
- **MLPRegressor** also supports **multi-output regression**, in which a sample can have more than one target.

Training data

Array X : holds the training samples



shape \rightarrow (n_samples, n_features)

Array y : holds the target



shape \rightarrow (n_samples,)

MLPClassifier

- How to implement MLPClassifier?

Step 1: Instantiate a **MLP** classifier estimator.

```
1 from sklearn.neural_network import MLPClassifier
2 MLP_clf = MLPClassifier()
```

Step 2: Call **fit** method on **MLP classifier object** with **training feature matrix** and **label vector** as arguments.

```
1 # Model training with feature matrix X_train and
2 # label vector or matrix y_train
3 MLP_clf.fit(X_train, y_train)
```

MLPClassifier

Step 3: After fitting (training), the model can make predictions for new samples (X_{test}) using two methods:

```
1 MLP_clf.predict(X_test)
2 MLP_clf.predict_proba(X_test)
```

predict →

- gives labels for new samples
- for example:
`array([1, 0])`

predict_proba →

- gives vector of probability estimates per sample
- for example:
`array([1.967...e-04, 9.998...-01])`
- MLPClassifier supports only the **Cross-Entropy loss function**

How to set the number of hidden layers?

`hidden_layer_sizes`


- This parameter sets the number of layers and the number of neurons in each layer.
- It is a **tuple** where each element in the tuple represents the number of neurons at the i th position where i is the index of the tuple.
- The length of tuple denotes the total number of hidden layers in the network.

To create a 3 hidden layer neural network with 15 neurons in first layer, 10 neurons in second layer and 5 neurons in third layer:

```
1 MLPClassifier(hidden_layer_sizes=(15,10,5))
```

How to perform **regularization** in MLPClassifier?

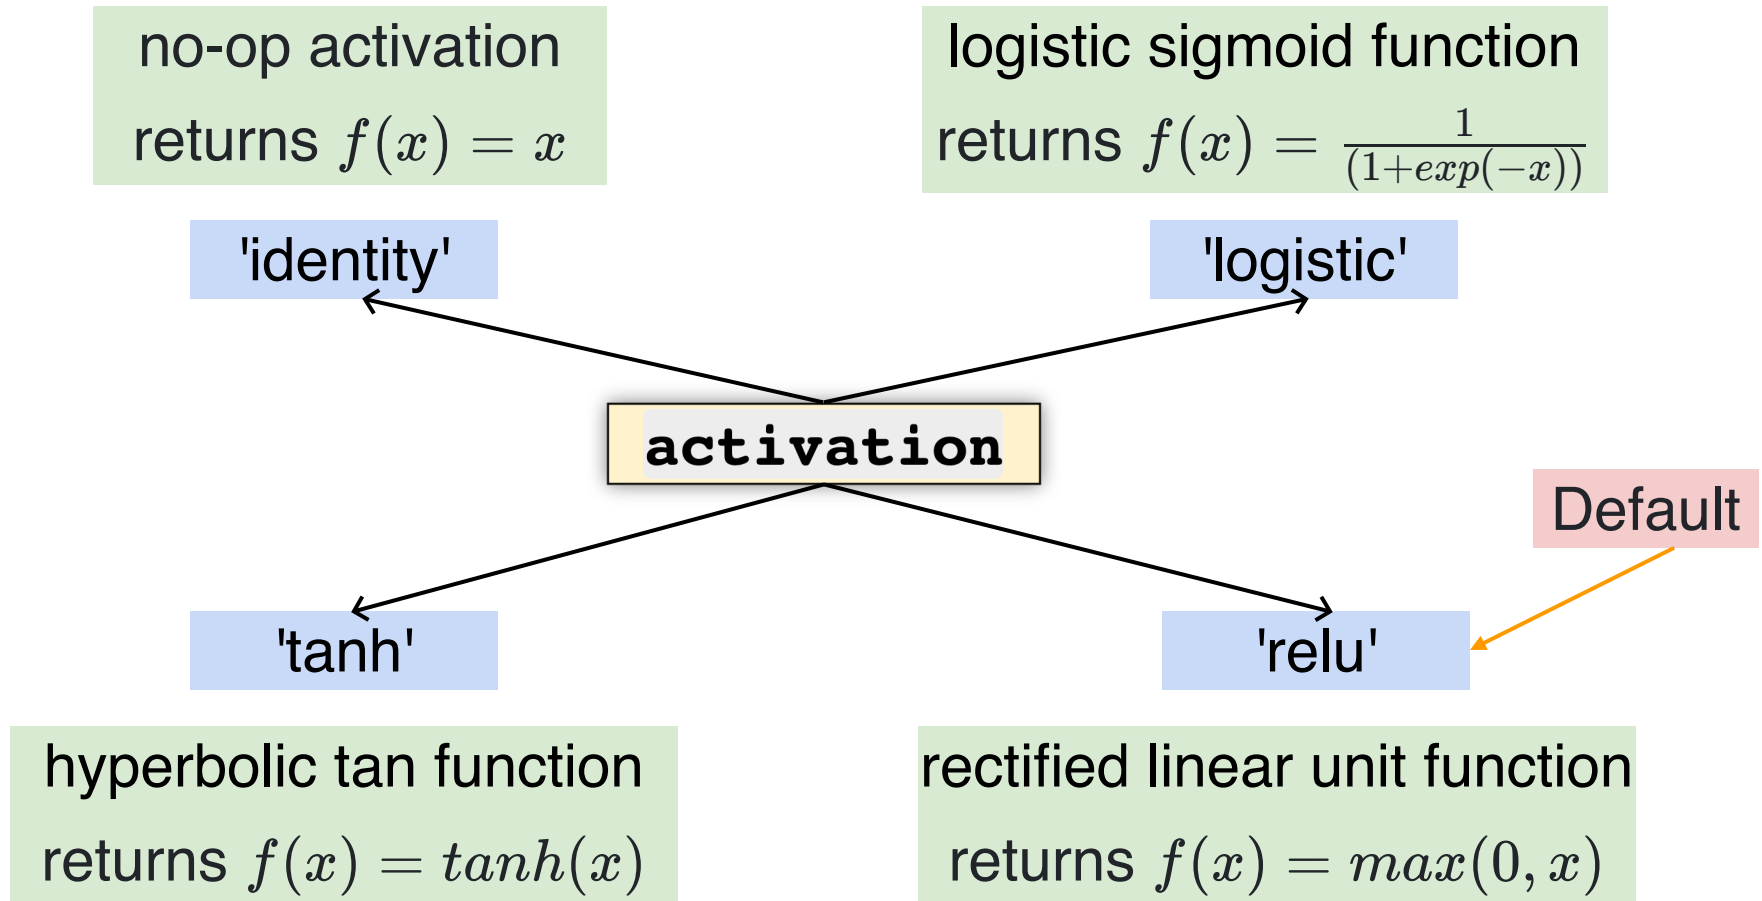
- The alpha parameter sets L2 penalty Regularization parameter



Default:


```
1 alpha = 0.0001
```


How to set the **activation function** for the hidden layers?



How to perform **weight optimization** in MLPClassifier?

- MLPClassifier optimizes the log-loss function using LBFGS or stochastic gradient descent



- If the **solver** is 'lbfgs', the classifier will not use minibatch.
- Size of minibatches can be set to other stochastic optimizers: **batch_size** (int)

- default batch_size is 'auto'.

↓

```
1 batch_size=min(200, n_samples)
```

How to view **weight matrix coefficients** of trained MLPClassifier?

coefs_

- It is a **list** of shape (n_layers - 1,)
- The i th element in the list represents the weight matrix corresponding to layer i .

Example:

- "weights between input and first hidden layer:"

```
1 print(MLP_clf.coefs_[0])
```

- "weights between first hidden and second hidden layer:"

```
1 print(MLP_clf.coefs_[1])
```

weights between input and first hidden layer:

```
[[-0.14203691 -1.18304359 -0.85567518 -4.53250719 -0.60466275]
 [-0.69781111 -3.5850093 -0.26436018 -4.39161248 0.06644423]]
```

weights between first hidden and second hidden layer:

```
[ [ 0.29179638 -0.14155284]
 [ 4.02666592 -0.61556475]
 [-0.51677234 0.51479708]
 [ 7.37215202 -0.31936965]
 [ 0.32920668 0.64428109]]
```

How to view **bias vector** of trained MLPClassifier?

intercepts_

- It is a **list** of shape $(n_layers - 1,)$
- The i th element in the list bias vector corresponding to layer $i + 1$.

Example:

- "Bias values for first hidden layer:"

```
1 print(MLP_clf.intercepts_[0])
```

- "Bias values for second hidden layer:"

```
1 print(MLP_clf.intercepts_[1])
```

```
Bias values for first hidden layer:  
[-0.14962269 -0.59232707 -0.5472481  7.02667699 -0.87510813]  
  
Bias values for second hidden layer:  
[-3.61417672 -0.76834882]
```

Some parameters in MLPClassifier

learning_rate

'constant'

'invscaling'

'adaptive'

default: 'constant'

learning_rate_init

float value

default: 0.001

power_t

float value

default: 0.5

max_iter

int value

default: 500

- **learning_rate** and **power_t** are used only for **solver = 'sgd'**
- **learning_rate_init** is used when **solver='sgd'** or 'adam'.
- **shuffle** is used to shuffle samples in each iteration when **solver='sgd'** or 'adam'
- **momentum** is used for gradient descent update when **solver='sgd'**

MLPRegressor

- MLPRegressor trains using backpropagation with no activation function in the output layer.
- Therefore, it uses the square error as the loss function, and the output is a set of continuous values.

The parameters of MLPRegressor are the same as that of MLPClassifier.

How to **implement** MLPRegressor?

Step 1: Instantiate a **MLP** regressor estimator.

```
1 from sklearn.neural_network import MLPRegressor
2 MLP_reg = MLPRegressor()
```

Step 2: Call **fit** method on **MLP regressor object** with **training feature matrix** and **label vector** as arguments.

```
1 # Model training with feature matrix X_train and
2 # label vector or matrix y_train
3 MLP_reg.fit(X_train, y_train)
```

Step 3: After fitting (training), the model can make predictions for new samples (X_{test}):

```
1 MLP_reg.predict(X_test)
```

- returns predicted values for new samples
- for example:
`array([-0.9..., -7.1...])`

```
1 MLP_reg.score(X_test,y_test)
```

- returns R^2 score
- for example:
`0.45678889`